


Templates, Image Pyramids, and Filter Banks




Computer Vision

Derek Hoiem, University of Illinois


Review

1. Match the spatial domain image to the Fourier magnitude image



B

A



D



E

C

Reminder


- HW 1 due in one week

Today's class

- Template matching
- Image Pyramids
- Filter banks and texture
- Denoising, Compression

Template matching

- Goal: find  in image
- Main challenge: What is a good similarity or distance measure between two patches?
 - Correlation
 - Zero-mean correlation
 - Sum Square Difference
 - Normalized Cross Correlation




Matching with filters

- Goal: find  in image
- Method 0: filter the image with eye patch


$$h[m,n] = \sum_{k,l} g[k,l] f[m+k, n+l]$$



f = image
g = filter



Input




Filtered Image

What went wrong?


Matching with filters

- Goal: find  in image
- Method 1: filter the image with zero-mean eye


$$h[m,n] = \sum_{k,l} (f[k,l] - \bar{f}) \underbrace{(g[m+k, n+l])}_{\text{mean of } f}$$



Input



Filtered Image (scaled)



Thresholded Image


True detections

False detections

Matching with filters

- Goal: find  in image
- Method 2: SSD


$$h[m,n] = \sum_{k,l} (g[k,l] - f[m+k, n+l])^2$$



Input



1 - $\text{sqrt}(\text{SSD})$



Thresholded Image

True detections

Matching with filters

Can SSD be implemented with linear filters?


$$h[m, n] = \sum_{k, l} (g[k, l] - f[m + k, n + l])^2$$

Matching with filters

- Goal: find  in image
- Method 2: SSD

What's the potential downside of SSD?

$$h[m,n] = \sum_{k,l} (g[k,l] - f[m+k, n+l])^2$$



Input



1- sqrt(SSD)

Matching with filters

- Goal: find  in image
- Method 3: Normalized cross-correlation

$$h[m,n] = \frac{\sum_{k,l} (g[k,l] - \bar{g})(f[m-k, n-l] - \bar{f}_{m,n})}{\left(\sum_{k,l} (g[k,l] - \bar{g})^2 \sum_{k,l} (f[m-k, n-l] - \bar{f}_{m,n})^2 \right)^{0.5}}$$


mean template
↓
 $\sum_{k,l} (g[k,l] - \bar{g})(f[m-k, n-l] - \bar{f}_{m,n})$

mean image patch
↓

Matlab: `normxcorr2(template, im)`

Matching with filters


- Goal: find  in image
- Method 3: Normalized cross-correlation



Input




Normalized X-Correlation



True detections
↓
↓
Thresholded Image

Matching with filters


- Goal: find  in image
- Method 3: Normalized cross-correlation



Input



Normalized X-Correlation



True detections

Q: What is the best method to use?


A: Depends

- Zero-mean filter: fastest but not a great matcher
- SSD: next fastest, sensitive to overall intensity
- Normalized cross-correlation: slowest, invariant to local average intensity and contrast


Q: What if we want to find larger or smaller eyes?

A: Image Pyramid

Review of Sampling



Gaussian pyramid



Source: Forsyth


Template Matching with Image Pyramids

Input: Image, Template

1. Match template at current scale
2. Downsample image
3. Repeat 1-2 until image is very small
4. Take responses above some threshold, perhaps with non-maxima suppression

Coarse-to-fine Image Registration


1. Compute Gaussian pyramid
2. Align with coarse pyramid
3. Successively align with finer pyramids
 - Search smaller range





Why is this faster?

Are we guaranteed to get the same result?

Laplacian filter




Laplacian pyramid



Source: Forsyth

Computing Gaussian/Laplacian Pyramid



Can we reconstruct the original
from the laplacian pyramid?

Hybrid Image



Hybrid Image in Laplacian Pyramid

High frequency → Low frequency




Image representation

- Pixels: great for spatial resolution, poor access to frequency
- Fourier transform: great for frequency, not for spatial info
- Pyramids/filter banks: balance between spatial and frequency information

Major uses of image pyramids



- Compression
- Object detection
 - Scale search
 - Features
- Detecting stable interest points
- Registration
 - Course-to-fine

Application: Representing Texture




Source: Forsyth

Texture and Material



Texture and Orientation



Texture and Scale




What is texture?

Regular or stochastic patterns caused by
bumps, grooves, and/or markings

How can we represent texture?

- Compute responses of blobs and edges at various orientations and scales




Overcomplete representation: filter banks



Code for filter banks: www.robots.ox.ac.uk/~vgg/research/texclass/filters.html

Filter banks


- Process image with each filter and keep responses (or squared/abs responses)









How can we represent texture?

- Measure responses of blobs and edges at various orientations and scales
- Idea 1: Record simple statistics (e.g., mean, std.) of absolute filter responses

Can you match the texture to the response?





Representing texture by mean abs response



Mean abs responses

Representing texture


- Idea 2: take vectors of filter responses at each pixel and cluster them, then take histograms (more on in coming weeks)



Compression

How is it that a 4MP image can be compressed to a few hundred KB without a noticeable change?

Lossy Image Compression (JPEG)



Block-based Discrete Cosine Transform (DCT)

Using DCT in JPEG

- The first coefficient $B(0,0)$ is the DC component, the average intensity
- The top-left coeffs represent low frequencies, the bottom right – high frequencies





Image compression using DCT

- Quantize
 - More coarsely for high frequencies (which also tend to have smaller values)
 - Many quantized high frequency values will be zero
- Encode
 - Can decode with inverse dct

Filter responses

$$G = \begin{bmatrix} -415.38 & -30.19 & -61.20 & 27.24 & 56.13 & -20.10 & -2.39 & 0.46 \\ 4.47 & -21.86 & -60.76 & 10.25 & 13.15 & -7.09 & -8.54 & 4.88 \\ -46.83 & 7.37 & 77.13 & -24.56 & -28.91 & 9.93 & 5.42 & -5.65 \\ -48.53 & 12.07 & 34.10 & -14.76 & -10.24 & 6.30 & 1.83 & 1.95 \\ 12.12 & -6.55 & -13.20 & -3.95 & -1.88 & 1.75 & -2.79 & 3.14 \\ -7.73 & 2.91 & 2.38 & -5.94 & -2.38 & 0.94 & 4.30 & 1.85 \\ -1.03 & 0.18 & 0.42 & -2.42 & -0.88 & -3.02 & 4.12 & -0.66 \\ -0.17 & 0.14 & -1.07 & -4.19 & -1.17 & -0.10 & 0.50 & 1.68 \end{bmatrix}$$

\xrightarrow{u}

$\downarrow v$

Quantization table

$$Q = \begin{bmatrix} 16 & 11 & 10 & 16 & 24 & 40 & 51 & 61 \\ 12 & 12 & 14 & 19 & 26 & 58 & 60 & 55 \\ 14 & 13 & 16 & 24 & 40 & 57 & 69 & 56 \\ 14 & 17 & 22 & 29 & 51 & 87 & 80 & 62 \\ 18 & 22 & 37 & 56 & 68 & 109 & 103 & 77 \\ 24 & 35 & 55 & 64 & 81 & 104 & 113 & 92 \\ 49 & 64 & 78 & 87 & 103 & 121 & 120 & 101 \\ 72 & 92 & 95 & 98 & 112 & 100 & 103 & 99 \end{bmatrix}$$

Quantized values

$$B = \begin{bmatrix} -26 & -3 & -6 & 2 & 2 & -1 & 0 & 0 \\ 0 & -2 & -4 & 1 & 1 & 0 & 0 & 0 \\ -3 & 1 & 5 & -1 & -1 & 0 & 0 & 0 \\ -3 & 1 & 2 & -1 & 0 & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$



JPEG Compression Summary


1. Convert image to YCrCb
2. Subsample color by factor of 2
 - People have bad resolution for color
3. Split into blocks (8x8, typically), subtract 128
4. For each block
 - a. Compute DCT coefficients
 - b. Coarsely quantize
 - Many high frequency components will become zero
 - c. Encode (e.g., with Huffman coding)

Lossless compression (PNG)


1. Predict that a pixel's value based on its upper-left neighborhood
2. Store difference of predicted and actual value
3. Pkzip it (DEFLATE algorithm)

	c	b	d
a	x		


Denoising



Additive Gaussian Noise




Gaussian
Filter




Reducing Gaussian noise


$\sigma=0.05$



$\sigma=0.1$



$\sigma=0.2$



no
smoothing



$\sigma=1$ pixel



$\sigma=2$ pixels



Smoothing with larger standard deviations suppresses noise, but also blurs the image

Reducing salt-and-pepper noise by Gaussian smoothing

3x3



5x5




7x7



Alternative idea: Median filtering

- A **median filter** operates over a window by selecting the median intensity in the window




- Is median filtering linear?

Median filter


- What advantage does median filtering have over Gaussian filtering?
 - Robustness to outliers

filters have width 5 :





Median filter

Salt-and-pepper noise



Median filtered



- MATLAB: `medfilt2(image, [h w])`

Median vs. Gaussian filtering

3x3



5x5



7x7



Gaussian

Median



Other non-linear filters

- Weighted median (pixels further from center count less)
- Clipped mean (average, ignoring few brightest and darkest pixels)
- Bilateral filtering (weight by spatial distance *and* intensity difference)



Bilateral filtering

Review of last three days

Review: Image filtering


$$g[\cdot, \cdot] \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

$f[.,.]$

$h[.,.]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0



$$h[m, n] = \sum_{k,l} f[k, l] g[m+k, n+l]$$

Credit: S. Seitz

Image filtering

 $f[.,.]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

 $h[.,.]$

0	10									

$$h[m, n] = \sum_{k,l} f[k, l] g[m+k, n+l]$$

$$g[\cdot, \cdot] \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

Credit: S. Seitz

Image filtering

 $f[.,.]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	90	0
0	0	0	90	90	90	90	90	90	0
0	0	0	90	90	90	90	90	90	0
0	0	0	90	0	90	90	90	90	0
0	0	0	90	90	90	90	90	90	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

 $h[.,.]$


$$h[m, n] = \sum_{k, l} f[k, l] g[m+k, n+l]$$

$$g[\cdot, \cdot] \frac{1}{9}$$


1	1	1
1	1	1
1	1	1

Filtering in spatial domain


1	0	-1
2	0	-2
1	0	-1




$$\begin{matrix} * \\ = \end{matrix}$$




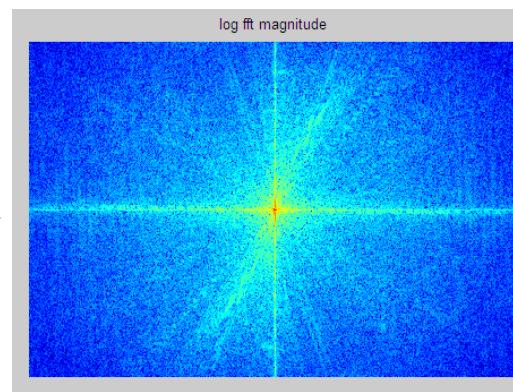
Filtering in frequency domain





FFT



FFT




Inverse FFT




Review of Last 3 Days

- Linear filters for basic processing
 - Edge filter (high-pass)
 - Gaussian filter (low-pass)


$$[-1 \ 1]$$



FFT of Gradient Filter




FFT of Gaussian



Gaussian

Review of Last 3 Days

- Derivative of Gaussian



Review of Last 3 Days

- Applications of filters
 - Template matching (SSD or Normxcorr2)
 - SSD can be done with linear filters, is sensitive to overall intensity
 - Gaussian pyramid
 - Coarse-to-fine search, multi-scale detection
 - Laplacian pyramid
 - More compact image representation
 - Can be used for compositing in graphics

Review of Last 3 Days

- Applications of filters
 - Downsampling
 - Need to sufficiently low-pass before downsampling
 - Compression
 - In JPEG, coarsely quantize high frequencies

Next class: edge and line detection

